

Exploring and comparing different merging algorithms for balanced binary search trees

Changhui (Eric) Zhou

October 9, 2025

IB Computer Science Extended Essay
Word Count: 3150 (Oct 8 version)

Contents

1	Introduction	2
2	Theory	2
2.1	Data structure terminology	2
2.2	Balanced binary search trees	3
2.3	Insertion-based merge	4
2.4	In-order traversal merge	5
2.5	Brown and Tarjan's merging algorithm (1979)	6
2.6	Optimality	8
3	Hypothesis	10
4	Experiment Design	10
4.1	Variables	10
4.1.1	Independent variables	10
4.1.2	Dependent variables	10
4.1.3	Controlled Variables	11
4.2	Procedure	11
5	Results	12
5.1	Raw data	12
5.2	Processed data	12
6	Discussion and Conclusion	16
6.1	Correlation between n and cost	16
6.2	Correlation between α and cost	16
7	Evaluation	17
	References	18

1 Introduction

A *data structure* is a way to store and organize data in order to facilitate access and modifications (Cormen, Leiserson, Rivest, & Stein, 2022). Designing and choosing more efficient data structures has always been a great pursuit for computer scientists, for optimal data structures can save huge amount of computing resources, especially in face of large amount of data. Basic data structures include ordered data structures like arrays, linked lists and binary search trees and unordered data structures like hashtables or disjoint set unions (DSU).

For ordered data structures, merging two or more instances of them while maintaining its ordered property may be frequently used in practice. For example, to investigate the factors affecting the school grade, data from different schools may be grouped and merged according to various factors. The efficiency of combination varies significantly based on the data structure itself and the algorithm used in the process. The merging of some data structures like linked lists, DSUs or heaps (especially leftist heaps) is simple and straightforward. However, these easy-to-merge data structures can lack of properties like random accessing (for linked lists or heaps) or ordered property (for DSUs). Balanced BSTs, however, supports a wide range of operations in just $O(\log n)$ time, and is widely used in theoretical researches, competitions and real life applications.

This essay will focus on investigating the theoretical time complexity (usually described using big-O notation, showing the asymptotic upper bound (Cormen et al., 2022)) and actual performance of merging algorithms of balanced BSTs, which are one of the most common and powerful data structure in real life.

Research question: How does different algorithm affect the efficiency of merging two instances of balanced search trees?

2 Theory

2.1 Data structure terminology

When a homogeneous relation (a binary relation between two elements) \leq on a set of elements X satisfies:

1. Antisymmetry: $\forall u, v \in X, (u \leq v \wedge v \leq u) \Leftrightarrow u = v.$
2. Totality: $\forall u, v \in X, u \leq v \vee v \leq u.$
3. Transitivity: $\forall u, v, w \in X, (u \leq v \wedge v \leq w) \Rightarrow u \leq w.$

We say $P = (X, \leq)$ is a total order. For example $P = (\mathbb{R}, \leq)$, where \leq is numerical comparison, is a total order. The set of finite strings and lexicographical order comparison is also a total order. But $P = (\{S : S \subset \mathbb{R}\}, \subset)$ (i.e. all the real number sets and the subset relation) is not a total order.

Ordered data structures can store elements that satisfies a total order while maintaining their order. In C++, arrays, vectors, linked lists and sets can be ordered data structures, but unordered sets (hashtables) are not ordered data structures.

Definition

Ordered data structures are data structures that can store elements that satisfies a total order while maintaining their order.

2.2 Balanced binary search trees

- A *graph* $G = (V, E)$ is the combination of the vertex set V and the edge set E .
- A *tree* $T = (V, E)$ is a connected acyclic graph.
 - The *root* of a tree is a designated vertex that has no parent. For every other vertex, exactly one of its neighbors is its parent. It can be easily proven that once the root is chosen, the tree is determined.
 - The *children* of a vertex are all its neighbors except for its parent. It can be easily proven that the parent of all the children of a vertex is the vertex itself.
 - The *leaf* of a tree is a vertex that has no children.

- The *subtree* of a vertex v is the union of v and the subtree of all its children, or only the vertex itself in the case that v is a leaf.
- The *depth* of a tree or a subtree is the maximum number of edges on the simple path from the root to one of the leaves.
- A *binary tree* is a tree that has no more than two children for each node.
- A *binary search tree (BST)* is a binary tree, whose nodes contain values under a total order, that has the following property: For any node v , all nodes in its left subtree are less than v , and all nodes in its right subtree are greater than v (Cormen et al., 2022).

Generally speaking, a balanced BST is a BST whose depth or the cost of iterating from the root to one specific leaf is strictly, expectedly or amortized $O(\log n)$. There are different kinds of balanced BSTs, like splay tree, treap, AVL trees, scapegoat trees and red-black trees.

In this essay, we will focus on AVL trees (named after Adelson-Velsky and Landis). AVL trees are a type of self-balancing binary search trees, which adjusts its shape through rotations and maintain the difference of the depths of two subtrees at most one (Karlton, Fuller, Scroggs, & Kaehler, 1976)¹. AVL trees are chosen because their properties make it easy to implement and analyze. AVL trees are known to have simple, straightforward balancing condition. Unlike splay or treap, each operation on an AVL tree have a predictable time complexity for every operation, as the shape of the tree is relatively static. All the rebalance operations happen locally, and no case-heavy analysis (like in red-black tree) or amortized operations (like in splay tree) are needed.

In this essay, we will assume that there are two instances of AVL trees T_1 and T_2 to be merged. Without losing generality, we will assume T_1 has n elements and T_2 has m elements and $n \geq m$.

2.3 Insertion-based merge

One of the basic operations supported by a balanced BST is insertion, where one element is added to the tree and the order of the tree is automatically

¹We assume the maximum depth difference is 1 in this essay, which might not be the case in early studies.

maintained. In fact, merging two instances of BSTs can be reduced to a sequence of insertions to a balanced BST. To be more specific, we iterate through all the elements in T_2 , insert them one by one into T_1 , that would be m operations with each having time complexity of $O(\log n)$, resulting in a overall complexity of $O(m \log n)$.

Algorithm: Insertion-Based Merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```

1: procedure INSERTIONBASEDMERGE( $T_1, T_2$ )
2:   for all elements  $x$  in  $T_2$  (in-order traversal) do
3:      $T_1$ . INSERT( $x$ )
4:   return  $T_1$ 

```

This algorithm performs well when m/n is small, as the overall time complexity will be mainly $O(\log n)$. However, when m and n are relatively at the same scale, the overall time complexity will be close to $O(n \log n)$.

2.4 In-order traversal merge

Another way to merge two instances of BSTs is to utilize the property that each instance is already in-order. To combine them, we can view this process as merging two sorted subarrays into a new array, just like a merge sort. The iteration and new array construction process take $O(n + m)$ time. With proper construction function, we can create a balanced BST in linear time out of a sorted array. Therefore, the overall time complexity is $O(n + m)$.

Algorithm: In-order traversal merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```
1: procedure MERGESORTBASEDMERGE( $T_1, T_2$ )
2:    $A_1 \leftarrow \text{INORDERTRAVERSAL}(T_1)$ 
3:    $A_2 \leftarrow \text{INORDERTRAVERSAL}(T_2)$ 
4:    $A \leftarrow \text{MERGESORTEDARRAYS}(A_1, A_2)$ 
5:    $T \leftarrow \text{BUILDBALANCEDBST}(A)$ 
6:   return  $T$ 
```

This algorithm performs well when m/n is large, as the overall time complexity will be approximately $O(n)$. However, when m is pretty negligible compared to n , a full iteration over T_1 will be still needed and the overall time complexity will still be $O(n)$, which wastes a lot of time.

In fact, it Stockmayer and Yao has proven that in term of number of comparisons, this algorithm is optimal when $m \leq n \leq \lfloor 3m/2 \rfloor + 1$ (Stockmeyer & Yao, 1980). This algorithm, however, does not perform well outside this range.

2.5 Brown and Tarjan's merging algorithm (1979)

In 1979, Brown and Tarjan proposed another algorithm based on the two merging algorithm mentioned above. It utilized both the tree structure for fast insertion-place location and the ordered property to reduce redundant operations. The algorithm again chooses the T_1 as the base tree and view the merging process as m insertions to a balanced BST of size n . However, the property that the inserted objects themselves are sorted helped to make the algorithm more efficient. Instead of iterating from the root, each insertion starts with the ending position of the last insertion, as it can be already told that the next insertion will happen to the right of the last insertion.

To be more specific, the algorithm keeps a stack called *path* and a stack called *successor*. The former is used to record the path from the root to the current node, while the latter records all the nodes on the *path* that is larger than the current node (that means they are on the right side of the current node their left children is visited on the *path*). Each insertion, instead of starting from the root, starts from the last node on the *successor* that is

smaller than the node to be inserted. Keep extending the *path* and *successor* during insertion. And the path shrinks back after the insertion, until a rebalance operation is triggered or we know that there need no rebalancing at all.

It is worth noticing that the rebalance operation may make the initial path unusable. In this case, we can simply dispose of the path under the rotated node start next insertion there (Brown & Tarjan, 1979).

Algorithm: Brown and Tarjan's Merging Algorithm

Require: Two balanced binary search trees T_1 (size n) and T_2 (size m), where $n \geq m$
Ensure: A single balanced BST containing all elements from T_1 and T_2

```

1: procedure FASTMERGE( $T_1, T_2$ )
2:   Initialize stack path  $\leftarrow \{\text{root}(T_1)\}$ 
3:   Initialize empty stack successor
4:   height  $\leftarrow \text{height}(T_1)$ 
5:   for all nodes  $x$  in  $T_2$  (in-order traversal) do
6:     Detach  $x$  from  $T_2$ 
7:      $\triangleright$  — Step 1: Adjust path to maintain PathPredicate —
8:     while successor not empty and  $\text{key}(x) >$ 
9:        $\text{key}(\text{top}(\text{successor}))$  do
10:       repeat
11:          $\text{pop}(\text{path})$ 
12:         until  $\text{top}(\text{path}) = \text{top}(\text{successor})$ 
13:          $\text{pop}(\text{successor})$ 
14:      $\triangleright$  — Step 2: Search down from last successor and insert  $x$  —
15:      $p \leftarrow \text{top}(\text{path})$ 
16:     while True do
17:       if  $\text{key}(x) < \text{key}(p)$  then
18:         if  $p.\text{left} = \text{Nil}$  then
19:            $p.\text{left} \leftarrow x$ ; break
20:         else
21:            $\text{push}(p, \text{successor}); p \leftarrow p.\text{left}$ 
22:       else
23:         if  $p.\text{right} = \text{Nil}$  then

```



```

21: | | | |  $p.right \leftarrow x$ ; break
22: | | | else
23: | | | |  $p \leftarrow p.right$ 
24: | | |  $\text{push}(p, \text{path})$ 
    | |  $\triangleright$  — Step 3: Adjust balance factors and rebalance if needed —
25: | | while  $\text{path}$  not empty do
26: | | |  $s \leftarrow \text{pop}(\text{path})$ 
27: | | | if tree at  $s$  is unbalanced then
28: | | | |  $\text{REBALANCE}(s)$ ; break
29: | | | else
30: | | | | Update balance factor of  $s$ 
31: | | | if  $\text{top}(\text{successor}) = s$  then
32: | | | |  $\text{pop}(\text{successor})$ 
33: | | return root of  $T_1$  (now merged)

```

2.6 Optimality

When merging two instances of size n and m respectively, there are in total $\binom{n+m}{n}$ possible outcomes. According to the decision tree theory, each of them corresponds to a decision tree leaf node. Since the merging algorithm is comparison based, the decision tree has to be a binary tree (i.e. Each node has at most two children). The height of the decision tree is therefore no lower than $O(\log_2(\binom{n+m}{n}))$.

According to Sterling's approximation,

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \quad (1)$$

which means

$$O(\log(n!)) = O\left(\frac{1}{2} \log(2\pi n) + n \log n - n \log e\right) = O(n \log n - n + O(\log n)) \quad (2)$$

Using the definition of combination number,

$$\binom{n+m}{m} = \frac{(n+m)!}{n!m!} \quad (3)$$

$$\log\left(\binom{n+m}{n}\right) = \log\left(\frac{(n+m)!}{n!}\right) \quad (4)$$

$$= (n+m) \log(n+m) - n \log n - m \log m + O(\log(n+m)) \quad (5)$$

$$= n \log\left(1 + \frac{m}{n}\right) + m \log\left(1 + \frac{n}{m}\right) + O(\log(n+m)) \quad (6)$$

Since $m \leq n$ we have $n \log(1 + m/n) \leq n \cdot (\frac{m}{n}) = m$, therefore the first term is $O(m)$. This means the first term should be neglected as $m \log(1 + \frac{n}{m})$ is the dominant term compared to $O(m)$.

Since $m \log(1 + \frac{n}{m})$ can be written as $m \log(n+m) - m \log(m)$, where the first term is more dominant than $O(\log(n+m))$, the third term should be neglected as well.

We can get the overall expression

$$\log\left(\binom{n+m}{n}\right) = O\left(m \log\left(1 + \frac{n}{m}\right)\right) \quad (7)$$

Theorem

The optimal time complexity of merging two instances of ordered data structures is $O(m \log(1 + \frac{n}{m}))$, multiplied by the comparison cost, which is assumed to be $O(1)$ in this case.

We can also consider two extreme cases, when $m \approx n$, and when $m \ll n$. In the first case, as m approaches n , the time complexity of the algorithm will become $O(n \log(1 + 1)) = O(n)$. This is significantly smaller than the insertion-based algorithm, which is $O(m \log(n)) \approx O(n \log n)$ and the same as the in-order traversal algorithm's $O(m + n)$, which is $O(n + n) = O(n)$ in this case. In the second case, as m approaches $O(1)$ the time complexity is $O(\log n)$, which is the same as the insertion-based algorithm and significantly more optimal than the in-order traversal algorithm.

The theorem and the extreme-case analysis both showed that Tarjan and Brown's algorithm is the optimal algorithm for merging two instances of ordered data structures.

3 Hypothesis

The cost for all the algorithms should increase as n increases. Insertion-based algorithm perform well when α is small while in-order transversal algorithm performs well when α is large. Brown and Tarjan's algorithm is the optimal algorithm for merging two instances of ordered data structures in general cases.

4 Experiment Design

4.1 Variables

4.1.1 Independent variables

The independent variables of this experiment are:

- $\alpha = \frac{n}{m} \in \{2^0, 2^3, 2^6, 2^9, 2^{12}\}$.
- $n \in \{2^{12}, 2^{13}, 2^{14}, 2^{15}, 2^{16}, 2^{17}, 2^{18}, 2^{19}, 2^{20}\}$.

It can be found that $m = \frac{n}{\alpha}$. However, n and m are not chosen as the independent variables because α can better represent the **relative scale** or **balance** of n and m , while n can better represent the **total scale** of data compared to m . We are more interested in these two properties, rather than the scale of one particular part of data.

α and n are both uniformly chosen in the log scale, since it can better help us investigate performance of the algorithm in varying data sizes. The change is expected to be insignificant if they were to be chosen in the linear scale.

4.1.2 Dependent variables

The dependent variable of the experiment is the efficiency of the algorithm. To be more specific, the efficiency is measured by the **clock time** taken by the algorithm t , as well as the **number of comparison operations** c .

These two variables can both represent the cost and measure the efficiency of the algorithm, but they have different foci and advantage over theoretical time-complexity analysis. The clock time t can reveal the invisible cost

neglected term during big-O analysis, better representing the real world efficiency, while the comparison operations c can better represent the actual number of comparisons, helping us to analyze the performance when the comparison is not actually $O(1)$ (e.g. lexicographical order can be more costly to compare.)

4.1.3 Controlled Variables

The test environment is controlled, as is listed in the Appendix

4.2 Procedure

The procedure of the experiment is as follows:

Testing Procedure

Require: Sets of parameter values A (for α) and N (for n).

Ensure: A CSV file `results.csv` containing time and comparison counts for each test.

```

1: procedure TEST( $A, N$ )
2:   Initialize random seed based on current time.
3:   Open results.csv and write header.
4:   for all Algorithm  $F$  in {"Insertion-based", "In-order traversal",
   "Brown and Tarjan's"} do
5:     for all  $\alpha \in A$  do
6:       for all  $n \in N$  do
7:         for trial  $\leftarrow 1$  to 20 do
8:           Randomly generate two AVL trees  $S_1$  and  $S_2$  of
           sizes  $n$  and  $n/\alpha$ .
9:           Record time  $t_1$ .
10:           $S_1 \leftarrow F(S_1, S_2)$ 
11:          Record time  $t_2$  and  $\Delta t \leftarrow t_2 - t_1$ .
12:          Record number of comparisons.
13:          Append ( $F, \alpha, n$ , trial,  $\Delta t$ , comparisons) to
           results.csv.
14:   Close results.csv.
```

5 Results

5.1 Raw data

The output is `seed = 524042979266800` and a 2701 row csv file. The output is too redundant to be shown here, and part of it is shown in Table 1.

The result is reproducible with the same seed since the random number generator (`std::mt19937`) is deterministic (`std::mersenne_twister_engine` — *cppreference.com*, 2024).

Table 1: Raw data (part of)

Method	α	N	Trial	Time (ms)	Comparisons
Insertion-based	1	4096	1	2.0169	90036
Insertion-based	1	4096	2	1.5414	91083
Insertion-based	1	4096	3	1.5598	91498
Insertion-based	1	4096	4	1.5160	90642
Insertion-based	1	4096	5	1.5823	89572
[...]					
Brown and Tarjan's	4096	1048576	16	0.5161	5079
Brown and Tarjan's	4096	1048576	17	0.5049	5175
Brown and Tarjan's	4096	1048576	18	0.5001	5088
Brown and Tarjan's	4096	1048576	19	0.4834	5147
Brown and Tarjan's	4096	1048576	20	0.5141	5180

5.2 Processed data

The processed data is also too redundant to be shown in this table, and part of it is shown in Table 2 and Table 3. To analyze the effect of α and n of the cost of the algorithm respectively, the cost is plotted against α and n , with the other parameters fixed, respectively in Figure 1 and Figure 3.

Table 2: Processed data for $\alpha = 64$ (time and comparisons vs. n).

n	Method	Time [ms]	Comparisons
4096	Insertion-based	0.022 ± 0.005	1154 ± 30
8192	Insertion-based	0.048 ± 0.013	2492 ± 42
[...]			
524288	Insertion-based	7.970 ± 1.105	233756 ± 486
1048576	Insertion-based	16.645 ± 2.057	492018 ± 375
4096	In-order traversal	0.369 ± 0.211	4121 ± 84
8192	In-order traversal	0.617 ± 0.198	8269 ± 123
[...]			
524288	In-order traversal	52.595 ± 5.530	532392 ± 128
1048576	In-order traversal	101.494 ± 6.763	1064745 ± 107
4096	Brown & Tarjan's	0.030 ± 0.003	699 ± 22
8192	Brown & Tarjan's	0.084 ± 0.081	1416 ± 40
16384	Brown & Tarjan's	0.133 ± 0.054	2830 ± 79
32768	Brown & Tarjan's	0.269 ± 0.061	5688 ± 92
65536	Brown & Tarjan's	0.645 ± 0.205	11382 ± 142
131072	Brown & Tarjan's	1.469 ± 0.553	22810 ± 145
262144	Brown & Tarjan's	3.519 ± 1.226	45675 ± 286
524288	Brown & Tarjan's	7.283 ± 0.823	91383 ± 384
1048576	Brown & Tarjan's	14.626 ± 1.152	182779 ± 555




Figure 1: Processed data (α fixed)

Table 3: Processed data for $n = 65536$ (time and comparisons vs. α).

α	Method	Time [ms]	Comparisons
1	Insertion-based	33.072 ± 3.972	1901789 ± 15394
8	Insertion-based	4.135 ± 0.998	207622 ± 4350
64	Insertion-based	0.712 ± 0.313	24630 ± 132
512	Insertion-based	0.111 ± 0.036	3073 ± 46
4096	Insertion-based	0.019 ± 0.007	386 ± 20
1	In-order traversal	12.241 ± 2.432	131068 ± 4
8	In-order traversal	6.462 ± 1.191	73722 ± 6
64	In-order traversal	5.985 ± 1.271	66486 ± 174
512	In-order traversal	5.822 ± 1.712	65059 ± 1384
4096	In-order traversal	5.536 ± 0.600	61662 ± 7352
1	Brown & Tarjan's	15.597 ± 3.442	256442 ± 246
8	Brown & Tarjan's	2.729 ± 0.518	56023 ± 258
64	Brown & Tarjan's	0.645 ± 0.205	11382 ± 142
512	Brown & Tarjan's	0.151 ± 0.080	1991 ± 50
4096	Brown & Tarjan's	0.027 ± 0.013	312 ± 15

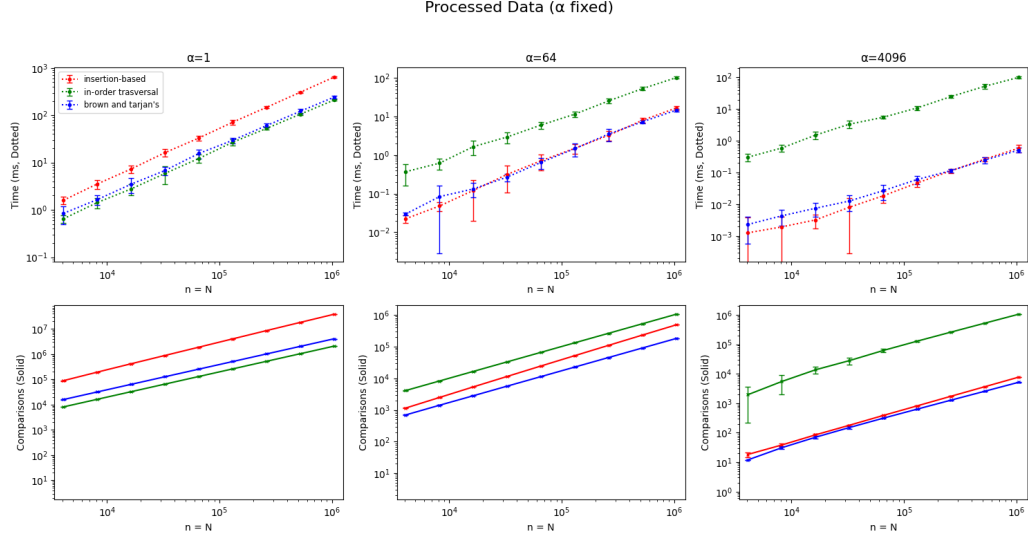


Figure 2: Processed data (α fixed) (both axis on log scale)

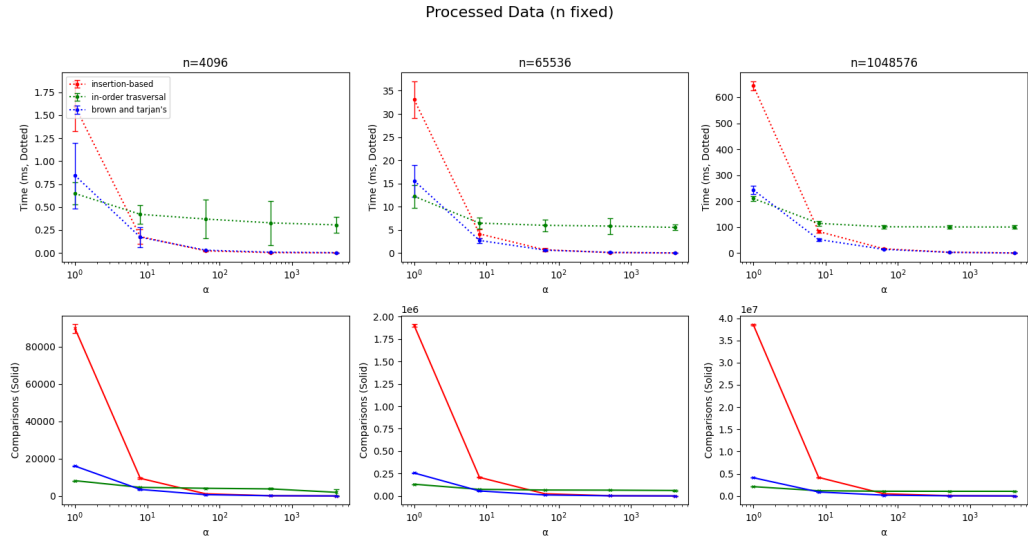


Figure 3: Processed data (n fixed)

6 Discussion and Conclusion

6.1 Correlation between n and cost

We can find clear trends from the graphs. In Figure 1, we can see that the cost of the algorithm all increases as the total data size n increases. It can be noticed that if we draw both axes in logarithmic scale, as is shown in Figure 2, the linear trend is more obvious. When α is fixed, the cost of three algorithms all increases linearly with the total data size n .

However, their rate of increase differs when α is different. For experiment where $\alpha = 1$, which means the sizes of both tree are the same, we can notice that both the time cost and the number of comparison of insertion-based algorithm is significantly higher than the other two algorithms, shown by a clear skyrocketing trend in the graph.

It can be also noticed that the despite having very significant disadvantage in number of comparison, the insertion based algorithm does not perform as bad as in the time. This might be because iteration through the smaller tree, which is obviously costly in this case as the tree takes up half the data, significantly costs time but not comparison.

On the contrary, for experiment where $\alpha = 64$, which means a small tree is being inserted to a significantly larger tree, the in-order traversal algorithm become the most costly, with both the time cost and number of comparisons skyrocketing. It can also be noticed that Brown and Tarjan's algorithm have smaller number of comparisons than the insertion-based algorithm, but similar cost in time. This might be because the cost of maintaining the path is quite significant. This implies if the cost of comparison is larger, like comparing the lexicographical order of the strings, Brown and Tarjan's algorithm's advantage will be more obvious. These trends are even more significant when $\alpha = 4096$.

There is also abnormal large error bar when both axes are drawn in log scale, which results from the exaggeration of the error bar at small values, since $\log x$ is more steep than x when x is small.

6.2 Correlation between α and cost

We can also find the advantage area of each algorithm in Figure 3. As is shown in the graph, insertion-based algorithm shows the least consistent

performance when α differs, with significant higher cost when $\alpha = 1$ and have lower cost when α is bigger. In-order transversal algorithm performs the best when α is small, but the cost reduces very slowly when α is bigger. This make it the algorithm with the most consistent performance. Brown and Tarjan’s algorithm performs slightly inferior to in-order transversal algorithm when α is small, and performs almost as well as the insertion-based algorithm when α is large. This made it the best algorithm overall. It can be also noticed that three algorithms’ performance is quite close when $\alpha \approx 8$.

The trend is almost the same throughout the graphs except for different scales, α is the only factor determining which algorithm is the optimal.

The error bars when n is small are significantly larger. The apparent instability for small n reflects the fact that the algorithm’s asymptotic behavior is masked by constant factors and system-level effects.

7 Evaluation

The experiment is conducted successfully with sufficient data and supports our hypothesis. Test data incorporated a large number of trials and a wide range of independent variables, making the experiment more robust. The results are reproducible and consistent with the hypothesis. However, there is still room for improvement.

Limitations and improvements:

- **Costly new and delete operations:** For the simplicity of implementation, `new` and `delete` operations are used to create and free nodes in the tree. Repetitive `new` and `delete` operations are expensive, risky and inconsistent with modern C++ standard (Meyers, 2014). A better implementation can be modern pointers and/or memory pool to avoid extra time cost.
- **Cache warmup and memory fragmentation control:** First several trials may be slower due to lack of cache warmup, later trials may be slower due to memory fragmentation. A better implementation is to run several warmup trials and use memory pool to avoid memory fragmentation.
- **Regular independent variables:** To avoid floating point calculation when calculating data size, the powers of 2 are used as the sizes of the

sequence generated. However, this results in a complete binary tree after tree construction, which is not always the case in practice. A better implementation is to use more independent variables and generate them randomly.

- **$O(1)$ comparison of int:** Only integer values whose comparison is $O(1)$ are put into the binary search trees. In many cases, the comparison can be more costly like in lexicographical order, floating point number comparison, big numbers or their ratios, etc. In these cases, the cost of logic or memory operations can be more insignificant compared to the cost of comparison, which means reducing the number of comparison is of greater priority. This can result in significantly different performance under the same data scale, as is discussed before. It would be better if more data types other than integer are used.
- **Fully shuffled dataset:** Only cases where the data is fully shuffled are considered. In many cases, the data is not purely random. Some algorithm may run well for random data but not every trial, like treap without random priority. When merging two instances of ordered data structures, elements in one instance may be significantly smaller than the other, in which case the performance can be different. Cases like this should be investigated as well.
- **Limited trials and independent variable range:** The experiment is conducted with a limited number of trials and a limited range of independent variables, due to the time limit and poor performance of the device. Some properties that can only be found with large n or certain dataset may not be investigated. It would be better if more trials and larger independent variable ranges are used with a better device.
- **No multithreading considered:** This essay did not discuss possible multithreading improvement of the algorithms.

References

Brown, M. R., & Tarjan, R. E. (1979). A fast merging algorithm. *Journal of the ACM (JACM)*, 26(2), 211–226.

- Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2022). *Introduction to algorithms* (Fourth ed.). MIT Press.
- Karltun, P. L., Fuller, S. H., Scroggs, R., & Kaehler, E. (1976). Performance of height-balanced trees. *Communications of the ACM*, 19(1), 23–28.
- Meyers, S. (2014). *Effective Modern C++: 42 specific ways to improve your use of C++11 and C++14*. Retrieved from <http://cds.cern.ch/record/1953418>
- `std::mersenne_twister_engine` — `cppreference.com`. (2024). https://en.cppreference.com/w/cpp/numeric/random/mersenne_twister_engine. (Last modified 21 August 2024)
- Stockmeyer, P. K., & Yao, F. F. (1980). On the optimality of linear merge. *SIAM Journal on Computing*, 9(1), 85–90.

Appendix A Test environment

Table 4: Test environment

Device:	Laptop
CPU:	Intel(R) Core(TM) Ultra 7 155H (1.40 GHz)
Memory:	32GB
OS:	Windows 11 24H2 26100.6584
Compiler (C++):	g++ 15.2.0 (algorithm implementation and timing)
Compiling Command:	<code>g++ test.cpp -o test -std=c++17 -O0 -Wall</code>
Interpreter:	Python 3.12.2 (graphing)

Appendix B Algorithm implementation

Listing 1: `avl_set.h`

```

1  #include <vector>
2  #include <algorithm>
3  #include <memory>
4  #include <list>
5  #include <functional>
6  #include <stack>
7  #ifdef DEBUG
8  #include <assert.h>
9  #endif
10
11 template <typename T, typename Compare = std::less<T>>
```

```

12 class AVLSet {
13 private:
14     struct Node {
15         T key;
16         Node* left;
17         Node* right;
18         int height;
19
20         template <typename... Args>
21         Node(Args&&... args)
22             : key(std::forward<Args>(args)...),
23               left(nullptr),
24               right(nullptr),
25               height(1) {}
26     };
27
28     Node* root;
29     size_t size;
30     Compare comp_;
31
32     // Helper functions for merging
33     int height(Node* node) const {
34         return node ? node->height : 0;
35     }
36
37     Node* create_node(const T& key) {
38         ++size;
39         return new Node(key);
40     }
41
42     Node* create_node(T&& key) {
43         ++size;
44         return new Node(std::move(key));
45     }
46
47     // Generates a balanced subtree in O(n) time out from a ordered sequence.
48     // Returns the root node pointer.
49     // *Preconditions
50     // keys have to be ordered
51     // _RandAccIt is the random access iterator
52     // (*bg) and (*ed) should be of type T
53
54     template<typename _RandAccIt>
55     Node* build(const _RandAccIt& bg, const _RandAccIt& ed){
56         if(bg == ed) return nullptr;
57         auto it = bg;
58         if(++it == ed) return create_node(*bg);
59         it = bg + (ed - bg) / 2; // The same as (bg + ed)/2 but avoids overflow problems
60         auto cur = create_node(*it);
61         cur->left = build(bg, it);
62         cur->right = build(it + 1, ed);
63         update_height(cur);
64         return cur;
65     }
66
67     void update_height(Node* node) {
68         node->height = 1 + std::max(height(node->left), height(node->right));

```

```

69     }
70
71     void delete_tree(Node* node) {
72         if (!node) return;
73         delete_tree(node->left);
74         delete_tree(node->right);
75         --size;
76         delete node;
77     }
78
79     Node* rotate_right(Node* y) {
80         Node* x = y->left;
81         Node* T2 = x->right;
82
83         x->right = y;
84         y->left = T2;
85
86         update_height(y);
87         update_height(x);
88
89         return x;
90     }
91
92     Node* rotate_left(Node* x) {
93         Node* y = x->right;
94         Node* T2 = y->left;
95
96         y->left = x;
97         x->right = T2;
98
99         update_height(x);
100        update_height(y);
101
102        return y;
103    }
104
105    int balance_factor(Node* node) const {
106        return node ? height(node->left) - height(node->right) : 0;
107    }
108
109    Node* balance(Node* node) {
110        update_height(node);
111        int bf = balance_factor(node);
112
113        // Left Heavy
114        if (bf > 1) {
115            if (balance_factor(node->left) < 0)
116                node->left = rotate_left(node->left);
117            return rotate_right(node);
118        }
119        // Right Heavy
120        if (bf < -1) {
121            if (balance_factor(node->right) > 0)
122                node->right = rotate_right(node->right);
123            return rotate_left(node);
124        }
125        return node;

```

```

126     }
127
128     template <typename Func>
129     void traverse_in_order(Node* node, Func f) const {
130         if (!node) return;
131         std::stack<Node*> stack;
132         Node* current = node;
133         while (current || !stack.empty()) {
134             while (current) {
135                 stack.push(current);
136                 current = current->left;
137             }
138             current = stack.top();
139             stack.pop();
140             f(current->key);
141             current = current->right;
142         }
143     }
144
145     Node* insert(Node* node, const T& key) {
146         if (!node) {
147             return create_node(key);
148         }
149
150         if (comp_(key, node->key)) {
151             node->left = insert(node->left, key);
152         } else if (comp_(node->key, key)) {
153             node->right = insert(node->right, key);
154         } else {
155             return node;
156         }
157
158         return balance(node);
159     }
160
161     public:
162     AVLSet(Compare Comp = Compare()) : root(nullptr), size(0), comp_(Comp) {}
163     Compare comparator() const { return comp_; }
164     // Move operations
165     AVLSet(AVLSet&& other) noexcept
166         : root(other.root),
167         size(other.size){
168         other.root = nullptr;
169         other.size = 0;
170     }
171
172     AVLSet& operator=(AVLSet&& other) noexcept {
173         if (this != &other) {
174             clear();
175             root = other.root;
176             size = other.size;
177             other.root = nullptr;
178             other.size = 0;
179         }
180         return *this;
181     }
182

```

```

183 // Disable copy operations
184 AVLSet(const AVLSet&) = delete;
185 AVLSet& operator=(const AVLSet&) = delete;
186
187 void clear() {
188     delete_tree(root);
189     root = nullptr;
190 }
191
192 bool empty() const {
193     return size == 0;
194 }
195
196 size_t get_size() const {
197     return size;
198 }
199
200 template <typename Func>
201 void traverse_in_order(Func f) const {
202     traverse_in_order(root, f);
203 }
204
205 std::vector<T> items() const {
206     std::vector<T> result;
207     traverse_in_order([&result](const T& key) {
208         result.push_back(key);
209     });
210     return result;
211 }
212
213
214 void swap_with(AVLSet& other) {
215     std::swap(root, other.root);
216     std::swap(size, other.size);
217 }
218
219 void insert(const T& val){
220     root = insert(root, val);
221 }
222 void remove(const T& val){
223     root = remove(root, val);
224 }
225 template<typename RandAccIt>
226 void construct(RandAccIt bg, RandAccIt ed){
227     // Clear the current tree
228     delete_tree(root);
229     root = nullptr;
230     // Build new tree
231     root = build(bg, ed);
232     size = static_cast<size_t>(ed - bg);
233 }
234
235
236 /** Merge two sets in O(N+M) time.
237  * Some additional space may be costed.
238  * But it does not affect the result of the experiment.
239  * @param other The AVLSet to be merged into this set.

```



```

240  */
241
242  void linearmerge(AVLSet&& other) {
243      if (other.empty()) return;
244      std::vector<T> q1 = items();
245      std::vector<T> q2 = other.items();
246      std::vector<T> all_elements;
247      all_elements.reserve(q1.size() + q2.size());
248      std::merge(q1.begin(), q1.end(), q2.begin(), q2.end(),
249                std::back_inserter(all_elements), comp_);
250      delete_tree(root);
251      delete_tree(other.root);
252      root = nullptr;
253      other.root = nullptr;
254      root = build(all_elements.begin(), all_elements.end());
255      size = all_elements.size();
256  }
257
258  /**
259   * Merge two sets in  $O(M \log(N))$  time.
260   * @param other The AVLSet to be merged into this set.
261   */
262  void simplemerge(AVLSet&& other) {
263      if (other.empty()) return;
264      if (size < other.size) { swap_with(other); }
265      other.traverse_in_order([this](const T& key) {
266          this->insert(key);
267      });
268      other.clear();
269  }
270
271  /**
272   * Merge two sets in  $O(M \log(1+N/M))$  time.
273   * @param other The AVLSet to be merged into this set.
274   */
275  void brownmerge(AVLSet&& other) {
276      if (other.empty()) return;
277      if (size < other.size) swap_with(other);
278
279      std::vector<T> elems = other.items();
280      other.clear();
281
282      // stacks of pointers-to-links (Node**). Each points to some parent->left or
283      // parent->right or &root
284      std::vector<Node**> path;
285      std::vector<Node**> successor;
286
287      path.push_back(&root);
288
289      for (const T& x : elems) {
290          while (!successor.empty() && !comp_(x, (*successor.back())->key)) {
291              Node** succLink = successor.back();
292              while (!path.empty() && path.back() != succLink) path.pop_back();
293              successor.pop_back();
294          }
295
296          Node** curLink = path.empty() ? &root : path.back();

```

```

296     Node* p = *curLink;
297     if (!p) {
298         *curLink = create_node(x);
299         path.push_back(curLink);
300     } else {
301         for (;;) {
302             path.push_back(curLink);
303             if (comp(x, p->key)) {
304                 if (p->left == nullptr) {
305                     p->left = create_node(x);
306                     path.push_back(&(p->left));
307                     break;
308                 } else {
309                     successor.push_back(curLink);
310                     curLink = &(p->left);
311                     p = *curLink;
312                 }
313             } else {
314                 if (p->right == nullptr) {
315                     p->right = create_node(x);
316                     path.push_back(&(p->right));
317                     break;
318                 } else {
319                     curLink = &(p->right);
320                     p = *curLink;
321                 }
322             }
323         }
324     }
325     while (!path.empty()) {
326         Node** link = path.back();
327         Node* s = *link;
328         path.pop_back();
329         if (!successor.empty() && successor.back() == link) successor.pop_back();
330         update_height(s);
331         int bf = balance_factor(s);
332         if (std::abs(bf) > 1) {
333             Node* newsub = balance(s);
334             *link = newsub;
335             while (!path.empty() && path.back() != link) path.pop_back();
336             break;
337         }
338
339         if (bf == 0) {
340             break;
341         }
342     }
343 }
344 #ifdef DEBUG
345 auto items_after = items();
346 for (size_t i = 1; i < items_after.size(); ++i)
347     assert(!comp(items_after[i], items_after[i-1]));
348 #endif
349 }
350
351 };

```

Listing 2: test.cpp

```

1  #include <iostream>
2  #include <vector>
3  #include <random>
4  #include <chrono>
5  #include <fstream>
6  #include "avl_set.h"
7
8  typedef std::shared_ptr<std::size_t> CounterPtr;
9
10 struct SharedCountingComp {
11     CounterPtr counter;
12
13     SharedCountingComp() : counter(std::make_shared<std::size_t>(0)) {}
14     SharedCountingComp(CounterPtr c) : counter(std::move(c)) {}
15
16     bool operator()(int a, int b) const {
17         ++(*counter);
18         return a < b;
19     }
20     size_t getcount() const { return *counter; }
21 };
22
23 const std::string methodnames[] = {"Insertion-based", "In-order traversal", "Brown and
    Tarjan's"};
24
25 std::tuple<double, size_t> testmerge(int type, std::mt19937_64& rnd, size_t q1 = 5e5,
    size_t q2 = 5e5) {
26     CounterPtr sharecounter = std::make_shared<std::size_t>(0);
27     auto comp = SharedCountingComp(sharecounter);
28     std::vector<int> vec(q1 + q2);
29     iota(vec.begin(), vec.end(), 0);
30     std::shuffle(vec.begin(), vec.end(), rnd);
31     auto s1 = AVLSet<int, SharedCountingComp>(comp), s2 = AVLSet<int, SharedCountingComp>
        >(comp);
32     sort(vec.begin()+0, vec.begin()+q1);
33     s1.construct(vec.begin() + 0, vec.begin() + q1);
34     sort(vec.begin() + q1 + 1, vec.begin() + q1 + q2);
35     s2.construct(vec.begin() + q1 + 1, vec.begin() + q1 + q2);
36     const auto t1 = std::chrono::steady_clock::now();
37     if(type == 1)
38         s1.simplemerge(std::move(s2));
39     else if (type == 2)
40         s1.linearmerge(std::move(s2));
41     else if (type == 3)
42         s1.brownmerge(std::move(s2));
43     const auto t2 = std::chrono::steady_clock::now();
44     double ms = std::chrono::duration<double, std::milli>(t2 - t1).count();
45     return std::make_tuple(ms, s1.comparator().getcount());
46 }
47
48 int main() {
49     auto seedx = std::chrono::steady_clock::now().time_since_epoch().count();
50     std::mt19937_64 rnd(seedx);
51
52     std::ofstream outfile("results.csv");

```

```

53     outfile << "method,alpha,N,trial,time_ms,comparisons\n"; // CSV header
54
55     for (int type : {1, 2, 3}) {
56         for (int alpha : {1, 1<<3, 1<<6, 1<<9, 1<<12}) {
57             for (long long n : {1<<12, 1<<13, 1<<14, 1<<15, 1<<16, 1<<17, 1<<18, 1<<19,
58                 1<<20}) {
59                 for (int trial = 1; trial <= 20; trial++) {
60                     auto [time, comp] = testmerge(type, rnd, n, n / alpha);
61                     outfile << methodnames[type - 1] << ", "
62                         << alpha << ", "
63                         << n << ", "
64                         << trial << ", "
65                         << time << ", "
66                         << comp << "\n";
67                 }
68             }
69         }
70
71     outfile.close();
72     std::cout << "seed = " << seedx << std::endl;
73     return 0;
74 }

```