

Exploring different merging algorithms for balanced trees and their time complexity optimization.

Changhui (Eric) Zhou

October 5, 2025

word count: ???

Contents

1	Introduction	2
2	Theory	2
2.1	Data structure terminology	2
2.2	Balanced binary search trees	3
2.3	Insertion-based merge	4
2.4	In-order traversal merge	4
2.5	Brown and Tarjan's merging algorithm (1979)	5
2.6	Optimality	7
3	Hypothesis	9
4	Experiment Design	9
4.1	Variables	9
4.1.1	Control Variables	9
	References	19

1 Introduction

A *data structure* is a way to store and organize data in order to facilitate access and modifications (Cormen, Leiserson, Rivest, & Stein, 2022). Designing and choosing more efficient data structures has always been a great pursuit for computer scientists, for optimal data structures can save huge amount of computing resources, especially in face of large amount of data. Basic data structures include ordered data structures like arrays, linked lists and binary search trees and unordered data structures like hashtables.

For ordered data structures, merging two or more instances of them while maintaining its ordered property may be frequently used in practice. For example, to investigate the factors affecting the school grade, data from different schools may be grouped and merged according to various factors. The efficiency of combination varies significantly based on the data structure itself and the algorithm used in the process.

This essay will focus on investigating the theoretical time complexity (need definitions aa) and actual performance of merging algorithms of different data structures, namely arrays, and BSTs, which are the most commonly used data structure in real life.

Research question: How does different algorithm affect the efficiency of merging two instances of balanced search trees?

2 Theory

2.1 Data structure terminology

When a homogeneous relation (a binary relation between two elements) \leq on a set of elements X satisfies:

1. Antisymmetry: $\forall u, v \in X, (u \leq v \wedge v \leq u) \Leftrightarrow u = v.$
2. Totality: $\forall u, v \in X, u \leq v \vee v \leq u.$
3. Transitivity: $\forall u, v, w \in X, (u \leq v \wedge v \leq w) \Rightarrow u \leq w.$

We say $P = (X, \leq)$ is a total order. For example $P = (\mathbb{R}, \leq)$, where \leq is numerical comparison, is a total order. The set of finite strings and

lexographical order comparison is also a total order. But $P = (\{S : S \subset \mathbb{R}\}, \subset)$ is not a total order.

Ordered data structures can store elements that satisfies a total order while maintaining their order. In C++, arrays, vectors, linked lists and sets can be ordered data structures, but unordered sets (hashtables) are not ordered data structures.

Definition

Ordered data structures are data structures that can store elements that satisfies a total order while maintaining their order.

2.2 Balanced binary search trees

- A *graph* $G = (V, E)$ is the combination of the vertex set V and the edge set E .
- A *tree* $T = (V, E)$ is a connected acyclic graph.
- A *binary tree* is a tree that has no more than two children for each node.
- A *binary search tree (BST)* is a binary tree, whose nodes contain values under a total order, that has the following property: For any node v , all nodes in its left subtree are less than v , and all nodes in its right subtree are greater than v (Cormen et al., 2022).

Generally speaking, a balanced BST is a BST whose depth or the cost of iterating from the root to one specific leaf is strictly, expectedly or amortized $O(\log n)$. There are different kinds of balanced BSTs, like splay tree, treap, AVL trees and red-black trees. This essay will focus on AVL trees. AVL trees are a type of self-balancing binary search trees, which adjusts its shape through rotations and maintain the difference of the depths of two subtrees at most one (Karlton, Fuller, Scroggs, & Kaehler, 1976)¹.

In this essay, we will assume that there are two instances of AVL trees T_1 and T_2 to be merged. Without losing generality, we will assume T_1 has n elements and T_2 has m elements and $n \geq m$.

¹In fact, this kind of BST was referred to as HB[1], but AVL trees nowadays are mainly HB[1]

2.3 Insertion-based merge

One of the basic operations supported by a balanced BST is insertion, where one element is added to the tree and the order of the tree is automatically maintained. In fact, merging two instances of BSTs can be reduced to a sequence of insertions to a balanced BST. To be more specific, we iterate through all the elements in T_2 , insert them one by one into T_1 , that would be m operations with each having time complexity of $O(\log n)$, resulting in an overall complexity of $O(m \log n)$.

Algorithm: Insertion-Based Merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```
1: procedure INSERTIONBASEDMERGE( $T_1, T_2$ )
2:   for all elements  $x$  in  $T_2$  (in-order traversal) do
3:      $T_1$ . INSERT( $x$ )
4:   return  $T_1$ 
```

This algorithm performs well when m/n is small, as the overall time complexity will be mainly $O(\log n)$. However, when m and n are relatively at the same scale, the overall time complexity will be close to $O(n \log n)$.

2.4 In-order traversal merge

Another way to merge two instances of BSTs is to utilize the property that each instance is already in-order. To combine them, we can view this process as merging two sorted subarrays into a new array, just like a merge sort. The iteration and new array construction process take $O(n + m)$ time. With proper construction function, we can create a balanced BST in linear time out of a sorted array. Therefore, the overall time complexity is $O(n + m)$.

Algorithm: Merge-Sort-Based BST Merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```
1: procedure MERGESORTBASEDMERGE( $T_1, T_2$ )
2:    $A_1 \leftarrow \text{INORDERTRAVERSAL}(T_1)$ 
3:    $A_2 \leftarrow \text{INORDERTRAVERSAL}(T_2)$ 
4:    $A \leftarrow \text{MERGESORTEDARRAYS}(A_1, A_2)$ 
5:    $T \leftarrow \text{BUILDBALANCEDBST}(A)$ 
6:   return  $T$ 
```

This algorithm performs well when m/n is large, as the overall time complexity will be approximately $O(n)$. However, when m is pretty negligible compared to n , a full iteration over T_1 will be still needed and the overall time complexity will still be $O(n)$, which wastes a lot of time.

In fact, it Stockmayer and Yao has proven that in term of number of comparisons, this algorithm is optimal when $m \leq n \leq \lfloor 3m/2 \rfloor + 1$ (Stockmeyer & Yao, 1980). This algorithm, however, does not perform well outside this range.

2.5 Brown and Tarjan's merging algorithm (1979)

In 1979, Brown and Tarjan proposed another algorithm based on the two merging algorithm mentioned above. It utilized both the tree structure for fast insertion-place location and the ordered property to reduce redundant operations. The algorithm again chooses the T_1 as the base tree and view the merging process as m insertions to a balanced BST of size n . However, the property that the inserted objects themselves are sorted helped to make the algorithm more efficient. Instead of iterating from the root, each insertion starts with the ending position of the last insertion, as it can be already told that the next insertion will happen to the right of the last insertion.

To be more specific, the algorithm keeps a stack called *path* and a stack called *successor*. The former is used to record the path from the root to the current node, while the latter records all the nodes on the *path* that is larger than the current node (that means they are on the right side of the current node their left children is visited on the *path*). Each insertion, instead of starting from the root, starts from the last node on the *successor* that is

smaller than the node to be inserted. Keep extending the *path* and *successor* during insertion. And the path shrinks back after the insertion, until a rebalance operation is triggered or we know that there need no rebalancing at all.

Algorithm: Brown-Tarjan Fast Height-Balanced Tree Merge

Require: Two balanced binary search trees T_1 (size n) and T_2 (size m), where $n \geq m$
Ensure: A single balanced BST containing all elements from T_1 and T_2

```

1: procedure FASTMERGE( $T_1, T_2$ )
2:   Initialize stack path  $\leftarrow \{\text{root}(T_1)\}$ 
3:   Initialize empty stack successor
4:   height  $\leftarrow \text{height}(T_1)$ 
5:   for all nodes  $x$  in  $T_2$  (in-order traversal) do
6:     Detach  $x$  from  $T_2$ 
7:      $\triangleright$  — Step 1: Adjust path to maintain PathPredicate —
8:     while successor not empty and  $\text{key}(x) >$ 
9:        $\text{key}(\text{top}(\text{successor}))$  do
10:      repeat
11:        |  $\text{pop}(\text{path})$ 
12:      until  $\text{top}(\text{path}) = \text{top}(\text{successor})$ 
13:       $\text{pop}(\text{successor})$ 
14:      $\triangleright$  — Step 2: Search down from last successor and insert  $x$  —
15:      $p \leftarrow \text{top}(\text{path})$ 
16:     while True do
17:       if  $\text{key}(x) < \text{key}(p)$  then
18:         if  $p.\text{left} = \text{Nil}$  then
19:           |  $p.\text{left} \leftarrow x$ ; break
20:         else
21:           |  $\text{push}(p, \text{successor}); p \leftarrow p.\text{left}$ 
22:       else
23:         if  $p.\text{right} = \text{Nil}$  then
24:           |  $p.\text{right} \leftarrow x$ ; break
25:         else
26:           |  $p \leftarrow p.\text{right}$ 

```

```

24: |   |   |   push( $p$ ,  $path$ )
    |   |   |   ▷ — Step 3: Adjust balance factors and rebalance if needed —
25: |   |   |   while  $path$  not empty do
26: |   |   |        $s \leftarrow \text{pop}(path)$ 
27: |   |   |       if tree at  $s$  is unbalanced then
28: |   |   |           REBALANCE( $s$ ); break
29: |   |   |       else
30: |   |   |           |   Update balance factor of  $s$ 
31: |   |   |           |   if  $\text{top}(successor) = s$  then
32: |   |   |           |       pop( $successor$ )
33: |   |   |   return root of  $T_1$  (now merged)

```

It is worth noticing that the rebalance operation may make the initial path unusable. In this case, we can simply dispose of the path under the rotated node start next insertion there (Brown & Tarjan, 1979).

2.6 Optimality

When merging two instances of size n and m respectively, there are in total $\binom{n+m}{n}$ possible outcomes. According to the decision tree theory, each of them corresponds to a decision tree leaf node. Since the merging algorithm is comparison based, the decision tree has to be a binary tree (i.e. Each node has at most two children). The height of the decision tree is therefore no lower than $O(\log_2(\binom{n+m}{n}))$.

According to Sterling's approximation,

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \quad (1)$$

which means

$$O(\log(n!)) = O\left(\frac{1}{2} \log(2\pi n) + n \log n - n \log e\right) = O(n \log n - n + O(\log n)) \quad (2)$$

Using the definition of combination number,

$$\binom{n+m}{m} = \frac{(n+m)!}{n!m!} \quad (3)$$

$$\log\left(\binom{n+m}{n}\right) = \log\left(\frac{(n+m)!}{n!}\right) \quad (4)$$

$$= (n+m) \log(n+m) - n \log n - m \log m + O(\log(n+m)) \quad (5)$$

$$= n \log\left(1 + \frac{m}{n}\right) + m \log\left(1 + \frac{n}{m}\right) + O(\log(n+m)) \quad (6)$$

Since $m \leq n$ we have $n \log(1 + m/n) \leq n \cdot (\frac{m}{n}) = m$, therefore the first term is $O(m)$. This means the first term should be neglected as $m \log(1 + \frac{n}{m})$ is the dominant term compared to $O(m)$.

Since $m \log(1 + \frac{n}{m})$ can be written as $m \log(n+m) - m \log(m)$, where the first term is more dominant than $O(\log(n+m))$, the third term should be neglected as well.

We can get the overall expression

$$\log\left(\binom{n+m}{n}\right) = O\left(m \log\left(1 + \frac{n}{m}\right)\right) \quad (7)$$

Theorem

The optimal time complexity of merging two instances of ordered data structures is $O(m \log(1 + \frac{n}{m}))$, multiplied by the comparison cost which is assumed to be $O(1)$ in this case.

3 Hypothesis

4 Experiment Design

4.1 Variables

4.1.1 Control Variables

Appendix

Test environment

- Device: Laptop
- CPU: Intel(R) Core(TM) Ultra 7 155H (1.40 GHz)
- Memory: 32GB
- OS: Windows 11 24H2 26100.6584
- Compiler (C++): g++ 15.2.0 (for algorithm efficiency test and timing)
- Interpreter: Python 3.12.2 (for batch test)

Listing 1: AvlSet

```
1 #include <vector>
2 #include <algorithm>
3 #include <memory>
4 #include <list>
5 #include <functional>
6 #include <stack>
7 #ifdef DEBUG
8 #include <assert.h>
9 #endif
10
11 template <typename T, typename Compare = std::less<T>>
12 class AVLSet {
13 private:
14     struct Node {
```

```

15     T key;
16     Node* left;
17     Node* right;
18     int height;
19
20     template <typename... Args>
21     Node(Args&&... args)
22         : key(std::forward<Args>(args)...),
23           left(nullptr),
24           right(nullptr),
25           height(1) {}
26 };
27
28 Node* root;
29 size_t size;
30 Compare comp_;
31
32 // Helper functions for merging
33 int height(Node* node) const {
34     return node ? node->height : 0;
35 }
36
37 Node* create_node(const T& key) {
38     ++size;
39     return new Node(key);
40 }
41
42 Node* create_node(T&& key) {
43     ++size;
44     return new Node(std::move(key));
45 }
46
47 // Generates a balanced subtree in O(n) time out from a
48 // ordered sequence.
49 // Returns the root node pointer.
50 // *Preconditions
51 // keys have to be ordered
52 // _RandAccIt is the random access iterator
53 // (*bg) and (*ed) should be of type T

```

```

54 template<typename _RandAccIt>
55 Node* build(const _RandAccIt& bg, const _RandAccIt& ed){
56     if(bg == ed) return nullptr;
57     auto it = bg;
58     if(++it == ed) return create_node(*bg);
59     it = bg + (ed - bg) / 2; // The same as (bg + ed)/2 but
        avoids overflow problems
60     auto cur = create_node(*it);
61     cur->left = build(bg, it);
62     cur->right = build(it + 1, ed);
63     update_height(cur);
64     return cur;
65 }
66
67 void update_height(Node* node) {
68     node->height = 1 + std::max(height(node->left), height(
        node->right));
69 }
70
71 void delete_tree(Node* node) {
72     if (!node) return;
73     delete_tree(node->left);
74     delete_tree(node->right);
75     --size;
76     delete node;
77 }
78
79 Node* rotate_right(Node* y) {
80     Node* x = y->left;
81     Node* T2 = x->right;
82
83     x->right = y;
84     y->left = T2;
85
86     update_height(y);
87     update_height(x);
88
89     return x;
90 }
91

```

```

92 Node* rotate_left(Node* x) {
93     Node* y = x->right;
94     Node* T2 = y->left;
95
96     y->left = x;
97     x->right = T2;
98
99     update_height(x);
100    update_height(y);
101
102    return y;
103 }
104
105 int balance_factor(Node* node) const {
106     return node ? height(node->left) - height(node->right) :
107         0;
108 }
109
110 Node* balance(Node* node) {
111     update_height(node);
112     int bf = balance_factor(node);
113
114     // Left Heavy
115     if (bf > 1) {
116         if (balance_factor(node->left) < 0)
117             node->left = rotate_left(node->left);
118         return rotate_right(node);
119     }
120     // Right Heavy
121     if (bf < -1) {
122         if (balance_factor(node->right) > 0)
123             node->right = rotate_right(node->right);
124         return rotate_left(node);
125     }
126     return node;
127 }
128
129 template <typename Func>
130 void traverse_in_order(Node* node, Func f) const {
131     if (!node) return;

```

```

131     std::stack<Node*> stack;
132     Node* current = node;
133     while (current || !stack.empty()) {
134         while (current) {
135             stack.push(current);
136             current = current->left;
137         }
138         current = stack.top();
139         stack.pop();
140         f(current->key);
141         current = current->right;
142     }
143 }
144
145 Node* insert(Node* node, const T& key) {
146     if (!node) {
147         return create_node(key);
148     }
149
150     if (comp_(key, node->key)) {
151         node->left = insert(node->left, key);
152     } else if (comp_(node->key, key)) {
153         node->right = insert(node->right, key);
154     } else {
155         return node;
156     }
157
158     return balance(node);
159 }
160
161 public:
162     AVLSet(Compare Comp = Compare()) : root(nullptr), size(0),
        comp_(Comp) {}
163     Compare comparator() const { return comp_; }
164     // Move operations
165     AVLSet(AVLSet&& other) noexcept
166         : root(other.root),
167         size(other.size){
168         other.root = nullptr;
169         other.size = 0;

```

```

170     }
171
172     AVLSet& operator=(AVLSet&& other) noexcept {
173         if (this != &other) {
174             clear();
175             root = other.root;
176             size = other.size;
177             other.root = nullptr;
178             other.size = 0;
179         }
180         return *this;
181     }
182
183     // Disable copy operations
184     AVLSet(const AVLSet&) = delete;
185     AVLSet& operator=(const AVLSet&) = delete;
186
187     void clear() {
188         delete_tree(root);
189         root = nullptr;
190     }
191
192     bool empty() const {
193         return size == 0;
194     }
195
196     size_t get_size() const {
197         return size;
198     }
199
200     template <typename Func>
201     void traverse_in_order(Func f) const {
202         traverse_in_order(root, f);
203     }
204
205     std::vector<T> items() const {
206         std::vector<T> result;
207         traverse_in_order([&result](const T& key) {
208             result.push_back(key);
209         });

```

```

210     return result;
211 }
212
213
214 void swap_with(AVLSet& other) {
215     std::swap(root, other.root);
216     std::swap(size, other.size);
217 }
218
219 void insert(const T& val){
220     root = insert(root, val);
221 }
222 void remove(const T& val){
223     root = remove(root, val);
224 }
225 template<typename RandAccIt>
226 void construct(RandAccIt bg, RandAccIt ed){
227     // Clear the current tree
228     delete_tree(root);
229     root = nullptr;
230     // Build new tree
231     root = build(bg, ed);
232     size = static_cast<size_t>(ed - bg);
233 }
234
235
236 /** Merge two sets in O(N+M) time.
237  * Some additional space may be costed.
238  * But it does not affect the result of the experiment.
239  * @param other The AVLSet to be merged into this set.
240  */
241
242 void linearmerge(AVLSet&& other) {
243     if (other.empty()) return;
244
245     // Get elements from both trees
246     std::vector<T> q1 = items();
247     std::vector<T> q2 = other.items();
248     std::vector<T> all_elements;
249     all_elements.reserve(q1.size() + q2.size());

```



```

250
251     // Merge sorted vectors
252     std::merge(q1.begin(), q1.end(), q2.begin(), q2.end(),
253               std::back_inserter(all_elements), comp_);
254
255     // Recycle both trees
256     delete_tree(root);
257     delete_tree(other.root);
258     root = nullptr;
259     other.root = nullptr;
260
261     // Build new tree
262     root = build(all_elements.begin(), all_elements.end());
263     size = all_elements.size();
264 }
265
266 /**
267  * Merge two sets in  $O(N \log(M))$  time.
268  * @param other The AVLSet to be merged into this set.
269  */
270 void simplemerge(AVLSet&& other) {
271     if (other.empty()) return;
272
273     if (size < other.size) { swap_with(other); }
274
275     // Insert all elements from the smaller tree (now 'other')
276     // into this
277     other.traverse_in_order([this](const T& key) {
278         this->insert(key);
279     });
280     other.clear();
281 }
282
283 void brownmerge(AVLSet&& other) {
284     if (other.empty()) return;
285     if (size < other.size) swap_with(other);
286
287     std::vector<T> elems = other.items();
288     other.clear();

```

```

289
290 // stacks of pointers-to-links (Node**). Each points to
291 // some parent->left or parent->right or &root
292 std::vector<Node**> path;
293 std::vector<Node**> successor;
294
295 path.push_back(&root);
296
297 for (const T& x : elems) {
298     // ==== CLIMB/RETRACT: pop successors while their key
299     // <= x (i.e. not (x < succ_key)) ====
300     while (!successor.empty() && !comp_(x, (*successor.
301         back())->key)) {
302         Node** succLink = successor.back();
303         // pop path entries until top == succLink
304         while (!path.empty() && path.back() != succLink)
305             path.pop_back();
306         successor.pop_back();
307     }
308
309     // ==== DESCEND from current finger (path.back()) to
310     // insertion point ====
311     Node** curLink = path.empty() ? &root : path.back();
312     Node* p = *curLink;
313     if (!p) {
314         // empty subtree (rare because root existed),
315         // insert directly
316         *curLink = create_node(x);
317         path.push_back(curLink);
318     } else {
319         for (;;) {
320             path.push_back(curLink); // link that points to
321                                     // p
322             if (comp_(x, p->key)) {
323                 // go left; mark this link as a successor (
324                 // we turned left here)
325                 if (p->left == nullptr) {
326                     p->left = create_node(x);
327                     path.push_back(&(p->left));
328                     break;

```

```

321         } else {
322             successor.push_back(curLink);
323             curLink = &(p->left);
324             p = *curLink;
325         }
326     } else {
327         // go right
328         if (p->right == nullptr) {
329             p->right = create_node(x);
330             path.push_back(&(p->right));
331             break;
332         } else {
333             curLink = &(p->right);
334             p = *curLink;
335         }
336     }
337 }
338 }
339
340 // ==== REBALANCE upward using the link-stack; attach
341 // rotated subtree via *link ====
342 while (!path.empty()) {
343     Node** link = path.back();
344     Node* s = *link;
345     path.pop_back();
346
347     if (!successor.empty() && successor.back() == link
348         ) successor.pop_back();
349
350     update_height(s);
351     int bf = balance_factor(s);
352
353     if (std::abs(bf) > 1) {
354         Node* newsub = balance(s); // returns new root
355                                     // of this subtree
356         *link = newsub; // reattach correctly via the
357                             // link
358
359         // retract path until the remaining entries are
360         // consistent with this rotation

```

```

356         while (!path.empty() && path.back() != link)
357             path.pop_back();
358         break; // stop climbing after performing
359             rotation
360     }
361
362     if (bf == 0) {
363         // height didn't increase => stop climbing
364         break;
365     }
366     // else continue climbing
367 }
368
369 #ifdef DEBUG
370 auto items_after = items();
371 for (size_t i = 1; i < items_after.size(); ++i)
372     assert(!comp_(items_after[i], items_after[i-1]));
373 #endif
374 }
375 };

```

References

- Brown, M. R., & Tarjan, R. E. (1979). A fast merging algorithm. *Journal of the ACM (JACM)*, 26(2), 211–226.
- Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2022). *Introduction to algorithms* (Fourth ed.). MIT Press.
- Karltun, P. L., Fuller, S. H., Scroggs, R., & Kaehler, E. (1976). Performance of height-balanced trees. *Communications of the ACM*, 19(1), 23–28.
- Stockmeyer, P. K., & Yao, F. F. (1980). On the optimality of linear merge. *SIAM Journal on Computing*, 9(1), 85–90.