

Exploring different merging algorithms for balanced trees and their time complexity optimization.

Changhui (Eric) Zhou

October 7, 2025

word count: ???

Contents

1	Introduction	2
2	Theory	2
2.1	Data structure terminology	2
2.2	Balanced binary search trees	3
2.3	Insertion-based merge	4
2.4	In-order traversal merge	4
2.5	Brown and Tarjan's merging algorithm (1979)	5
2.6	Optimality	7
3	Hypothesis	8
4	Experiment Design	8
4.1	Variables	8
4.1.1	Independent variables	8
4.1.2	Dependent variables	9
4.1.3	Controlled Variables	9
5	Results	9
5.1	Raw data	9
	References	17

1 Introduction

A *data structure* is a way to store and organize data in order to facilitate access and modifications (Cormen, Leiserson, Rivest, & Stein, 2022). Designing and choosing more efficient data structures has always been a great pursuit for computer scientists, for optimal data structures can save huge amount of computing resources, especially in face of large amount of data. Basic data structures include ordered data structures like arrays, linked lists and binary search trees and unordered data structures like hashtables.

For ordered data structures, merging two or more instances of them while maintaining its ordered property may be frequently used in practice. For example, to investigate the factors affecting the school grade, data from different schools may be grouped and merged according to various factors. The efficiency of combination varies significantly based on the data structure itself and the algorithm used in the process.

This essay will focus on investigating the theoretical time complexity (need definitions aa) and actual performance of merging algorithms of different data structures, namely arrays, and BSTs, which are the most commonly used data structure in real life.

Research question: How does different algorithm affect the efficiency of merging two instances of balanced search trees?

2 Theory

2.1 Data structure terminology

When a homogeneous relation (a binary relation between two elements) \leq on a set of elements X satisfies:

1. Antisymmetry: $\forall u, v \in X, (u \leq v \wedge v \leq u) \Leftrightarrow u = v.$
2. Totality: $\forall u, v \in X, u \leq v \vee v \leq u.$
3. Transitivity: $\forall u, v, w \in X, (u \leq v \wedge v \leq w) \Rightarrow u \leq w.$

We say $P = (X, \leq)$ is a total order. For example $P = (\mathbb{R}, \leq)$, where \leq is numerical comparison, is a total order. The set of finite strings and

lexographical order comparison is also a total order. But $P = (\{S : S \subset \mathbb{R}\}, \subset)$ is not a total order.

Ordered data structures can store elements that satisfies a total order while maintaining their order. In C++, arrays, vectors, linked lists and sets can be ordered data structures, but unordered sets (hashtables) are not ordered data structures.

Definition

Ordered data structures are data structures that can store elements that satisfies a total order while maintaining their order.

2.2 Balanced binary search trees

- A *graph* $G = (V, E)$ is the combination of the vertex set V and the edge set E .
- A *tree* $T = (V, E)$ is a connected acyclic graph.
- A *binary tree* is a tree that has no more than two children for each node.
- A *binary search tree (BST)* is a binary tree, whose nodes contain values under a total order, that has the following property: For any node v , all nodes in its left subtree are less than v , and all nodes in its right subtree are greater than v (Cormen et al., 2022).

Generally speaking, a balanced BST is a BST whose depth or the cost of iterating from the root to one specific leaf is strictly, expectedly or amortized $O(\log n)$. There are different kinds of balanced BSTs, like splay tree, treap, AVL trees and red-black trees. This essay will focus on AVL trees. AVL trees are a type of self-balancing binary search trees, which adjusts its shape through rotations and maintain the difference of the depths of two subtrees at most one (Karlton, Fuller, Scroggs, & Kaehler, 1976)¹.

In this essay, we will assume that there are two instances of AVL trees T_1 and T_2 to be merged. Without losing generality, we will assume T_1 has n elements and T_2 has m elements and $n \geq m$.

¹In fact, this kind of BST was referred to as HB[1], but AVL trees nowadays are mainly HB[1]

2.3 Insertion-based merge

One of the basic operations supported by a balanced BST is insertion, where one element is added to the tree and the order of the tree is automatically maintained. In fact, merging two instances of BSTs can be reduced to a sequence of insertions to a balanced BST. To be more specific, we iterate through all the elements in T_2 , insert them one by one into T_1 , that would be m operations with each having time complexity of $O(\log n)$, resulting in an overall complexity of $O(m \log n)$.

Algorithm: Insertion-Based Merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```
1: procedure INSERTIONBASEDMERGE( $T_1, T_2$ )  
2:   for all elements  $x$  in  $T_2$  (in-order traversal) do  
3:      $T_1$ . INSERT( $x$ )  
4:   return  $T_1$ 
```

This algorithm performs well when m/n is small, as the overall time complexity will be mainly $O(\log n)$. However, when m and n are relatively at the same scale, the overall time complexity will be close to $O(n \log n)$.

2.4 In-order traversal merge

Another way to merge two instances of BSTs is to utilize the property that each instance is already in-order. To combine them, we can view this process as merging two sorted subarrays into a new array, just like a merge sort. The iteration and new array construction process take $O(n + m)$ time. With proper construction function, we can create a balanced BST in linear time out of a sorted array. Therefore, the overall time complexity is $O(n + m)$.

Algorithm: In-order traversal merge

Require: Two balanced binary search trees T_1 and T_2

Ensure: A single balanced binary search tree containing all elements from T_1 and T_2

```
1: procedure MERGESORTBASEDMERGE( $T_1, T_2$ )
2:    $A_1 \leftarrow \text{INORDERTRAVERSAL}(T_1)$ 
3:    $A_2 \leftarrow \text{INORDERTRAVERSAL}(T_2)$ 
4:    $A \leftarrow \text{MERGESORTEDARRAYS}(A_1, A_2)$ 
5:    $T \leftarrow \text{BUILDBALANCEDBST}(A)$ 
6:   return  $T$ 
```

This algorithm performs well when m/n is large, as the overall time complexity will be approximately $O(n)$. However, when m is pretty negligible compared to n , a full iteration over T_1 will be still needed and the overall time complexity will still be $O(n)$, which wastes a lot of time.

In fact, it Stockmayer and Yao has proven that in term of number of comparisons, this algorithm is optimal when $m \leq n \leq \lfloor 3m/2 \rfloor + 1$ (Stockmeyer & Yao, 1980). This algorithm, however, does not perform well outside this range.

2.5 Brown and Tarjan's merging algorithm (1979)

In 1979, Brown and Tarjan proposed another algorithm based on the two merging algorithm mentioned above. It utilized both the tree structure for fast insertion-place location and the ordered property to reduce redundant operations. The algorithm again chooses the T_1 as the base tree and view the merging process as m insertions to a balanced BST of size n . However, the property that the inserted objects themselves are sorted helped to make the algorithm more efficient. Instead of iterating from the root, each insertion starts with the ending position of the last insertion, as it can be already told that the next insertion will happen to the right of the last insertion.

To be more specific, the algorithm keeps a stack called *path* and a stack called *successor*. The former is used to record the path from the root to the current node, while the latter records all the nodes on the *path* that is larger than the current node (that means they are on the right side of the current node their left children is visited on the *path*). Each insertion, instead of starting from the root, starts from the last node on the *successor* that is

smaller than the node to be inserted. Keep extending the *path* and *successor* during insertion. And the path shrinks back after the insertion, until a rebalance operation is triggered or we know that there need no rebalancing at all.

It is worth noticing that the rebalance operation may make the initial path unusable. In this case, we can simply dispose of the path under the rotated node start next insertion there (Brown & Tarjan, 1979).

Algorithm: Brown and Tarjan's Merging Algorithm

Require: Two balanced binary search trees T_1 (size n) and T_2 (size m), where $n \geq m$
Ensure: A single balanced BST containing all elements from T_1 and T_2

```

1: procedure FASTMERGE( $T_1, T_2$ )
2:   Initialize stack  $path \leftarrow \{\text{root}(T_1)\}$ 
3:   Initialize empty stack  $successor$ 
4:    $height \leftarrow \text{height}(T_1)$ 
5:   for all nodes  $x$  in  $T_2$  (in-order traversal) do
6:     Detach  $x$  from  $T_2$ 
7:      $\triangleright$  — Step 1: Adjust path to maintain PathPredicate —
8:     while  $successor$  not empty and  $\text{key}(x) >$ 
9:        $\text{key}(\text{top}(successor))$  do
10:      repeat
11:         $\text{pop}(path)$ 
12:      until  $\text{top}(path) = \text{top}(successor)$ 
13:       $\text{pop}(successor)$ 
14:      $\triangleright$  — Step 2: Search down from last successor and insert  $x$  —
15:      $p \leftarrow \text{top}(path)$ 
16:     while True do
17:       if  $\text{key}(x) < \text{key}(p)$  then
18:         if  $p.\text{left} = \text{Nil}$  then
19:            $p.\text{left} \leftarrow x$ ; break
20:         else
21:            $\text{push}(p, successor)$ ;  $p \leftarrow p.\text{left}$ 
22:       else
23:         if  $p.\text{right} = \text{Nil}$  then

```

```

21: | | | |  $p.right \leftarrow x$ ; break
22: | | | else
23: | | | |  $p \leftarrow p.right$ 
24: | | |  $\text{push}(p, \text{path})$ 
    | |  $\triangleright$  — Step 3: Adjust balance factors and rebalance if needed —
25: | | while  $\text{path}$  not empty do
26: | | |  $s \leftarrow \text{pop}(\text{path})$ 
27: | | | if tree at  $s$  is unbalanced then
28: | | | |  $\text{REBALANCE}(s)$ ; break
29: | | | else
30: | | | | Update balance factor of  $s$ 
31: | | | if  $\text{top}(\text{successor}) = s$  then
32: | | | |  $\text{pop}(\text{successor})$ 
33: | | return root of  $T_1$  (now merged)

```

2.6 Optimality

When merging two instances of size n and m respectively, there are in total $\binom{n+m}{n}$ possible outcomes. According to the decision tree theory, each of them corresponds to a decision tree leaf node. Since the merging algorithm is comparison based, the decision tree has to be a binary tree (i.e. Each node has at most two children). The height of the decision tree is therefore no lower than $O(\log_2(\binom{n+m}{n}))$.

According to Sterling's approximation,

$$n! \approx \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \quad (1)$$

which means

$$O(\log(n!)) = O\left(\frac{1}{2} \log(2\pi n) + n \log n - n \log e\right) = O(n \log n - n + O(\log n)) \quad (2)$$

Using the definition of combination number,

$$\binom{n+m}{m} = \frac{(n+m)!}{n!m!} \quad (3)$$

$$\log\left(\binom{n+m}{n}\right) = \log\left(\frac{(n+m)!}{n!}\right) \quad (4)$$

$$= (n+m) \log(n+m) - n \log n - m \log m + O(\log(n+m)) \quad (5)$$

$$= n \log\left(1 + \frac{m}{n}\right) + m \log\left(1 + \frac{n}{m}\right) + O(\log(n+m)) \quad (6)$$

Since $m \leq n$ we have $n \log(1 + m/n) \leq n \cdot (\frac{m}{n}) = m$, therefore the first term is $O(m)$. This means the first term should be neglected as $m \log(1 + \frac{n}{m})$ is the dominant term compared to $O(m)$.

Since $m \log(1 + \frac{n}{m})$ can be written as $m \log(n+m) - m \log(m)$, where the first term is more dominant than $O(\log(n+m))$, the third term should be neglected as well.

We can get the overall expression

$$\boxed{\log\left(\binom{n+m}{n}\right) = O(m \log(1 + \frac{n}{m}))} \quad (7)$$

Theorem

The optimal time complexity of merging two instances of ordered data structures is $O(m \log(1 + \frac{n}{m}))$, multiplied by the comparison cost with is assumed to be $O(1)$ in this case.

3 Hypothesis

4 Experiment Design

4.1 Variables

4.1.1 Independent variables

The independent variables of this experiment are:

- $\alpha = n/m \in \{1.0, 9.0, 99.0, 999.0\}$.
- $N = n + m \in \{10^4, 10^5, 10^6, 10^7, 10^8\}$

From the definition equation, we can find property shown in Equation 8,

$$n = \frac{\alpha N}{1 + \alpha}, m = \frac{N}{1 + \alpha} \quad (8)$$

which means for each pair of α and N , there is a corresponding pair of n and m . However, n and m are not chosen as the independent variables because α can better represent the **relative scale** or **balance** of n and m , while N can better represent the **total scale** of data. We are more interested in these two properties, rather than the scale of one particular part of data.

α and N are both almost uniformly chosen in the logarithmic scale, since it can better help us investigate the varying data size in application. The reason why α are not chosen in $10^k (k \in \mathbb{Z}^+)$ but $10^k - 1 (k \in \mathbb{Z}^+)$ is to avoid floating point calculation and error when choosing n and m .

4.1.2 Dependent variables

The dependent variable of the experiment is the efficiency of the algorithm. To be more specific, the efficiency is measured by the **clock time** taken by the algorithm t , as well as the **number of comparison operations** c .

These two variables can both represent the cost and measure the efficiency of the algorithm, but they have different foci and advantage over theoretical time-complexity analysis. The clock time t can reveal the invisible cost neglected term during big-O analysis, better representing the real world efficiency, while the comparison operations c can better represent the actual number of comparisons, helping us to analyze the performance when the comparison is not actually $O(1)$ (e.g. lexicographical order can be more costly to compare.)

4.1.3 Controlled Variables

5 Results

5.1 Raw data

The output is `seed = 524042979266800` and a 2701 row csv file, part of which is shown in Table 1.

Table 1: Raw data

Method	α	N	Trial	Time (ms)	Comparisons
Insertion-based	1	4096	1	2.0169	90036
Insertion-based	1	4096	2	1.5414	91083
Insertion-based	1	4096	3	1.5598	91498
Insertion-based	1	4096	4	1.5160	90642
Insertion-based	1	4096	5	1.5823	89572
[...]					
Brown and Tarjan's	4096	1048576	16	0.5161	5079
Brown and Tarjan's	4096	1048576	17	0.5049	5175
Brown and Tarjan's	4096	1048576	18	0.5001	5088
Brown and Tarjan's	4096	1048576	19	0.4834	5147
Brown and Tarjan's	4096	1048576	20	0.5141	5180

Appendix A Test environment

Table 2: Test environment

Device:	Laptop
CPU:	Intel(R) Core(TM) Ultra 7 155H (1.40 GHz)
Memory:	32GB
OS:	Windows 11 24H2 26100.6584
Compiler (C++):	g++ 15.2.0 (algorithm implementation and timing)
Interpreter:	Python 3.12.2 (graphing)

Appendix B Algorithm implementation

Listing 1: AvlSet

```

1  #include <vector>
2  #include <algorithm>
3  #include <memory>
4  #include <list>
5  #include <functional>
6  #include <stack>
7  #ifdef DEBUG
8  #include <assert.h>
9  #endif
10
11 template <typename T, typename Compare = std::less<T>>

```

```

12 class AVLSet {
13 private:
14     struct Node {
15         T key;
16         Node* left;
17         Node* right;
18         int height;
19
20         template <typename... Args>
21         Node(Args&&... args)
22             : key(std::forward<Args>(args)...),
23               left(nullptr),
24               right(nullptr),
25               height(1) {}
26     };
27
28     Node* root;
29     size_t size;
30     Compare comp_;
31
32     // Helper functions for merging
33     int height(Node* node) const {
34         return node ? node->height : 0;
35     }
36
37     Node* create_node(const T& key) {
38         ++size;
39         return new Node(key);
40     }
41
42     Node* create_node(T&& key) {
43         ++size;
44         return new Node(std::move(key));
45     }
46
47     // Generates a balanced subtree in O(n) time out from a ordered sequence.
48     // Returns the root node pointer.
49     // *Preconditions
50     // keys have to be ordered
51     // _RandAccIt is the random access iterator
52     // (*bg) and (*ed) should be of type T
53
54     template<typename _RandAccIt>
55     Node* build(const _RandAccIt& bg, const _RandAccIt& ed){
56         if(bg == ed) return nullptr;
57         auto it = bg;
58         if(++it == ed) return create_node(*bg);
59         it = bg + (ed - bg) / 2; // The same as (bg + ed)/2 but avoids overflow problems
60         auto cur = create_node(*it);
61         cur->left = build(bg, it);
62         cur->right = build(it + 1, ed);
63         update_height(cur);
64         return cur;
65     }
66
67     void update_height(Node* node) {
68         node->height = 1 + std::max(height(node->left), height(node->right));

```

```

69     }
70
71     void delete_tree(Node* node) {
72         if (!node) return;
73         delete_tree(node->left);
74         delete_tree(node->right);
75         --size;
76         delete node;
77     }
78
79     Node* rotate_right(Node* y) {
80         Node* x = y->left;
81         Node* T2 = x->right;
82
83         x->right = y;
84         y->left = T2;
85
86         update_height(y);
87         update_height(x);
88
89         return x;
90     }
91
92     Node* rotate_left(Node* x) {
93         Node* y = x->right;
94         Node* T2 = y->left;
95
96         y->left = x;
97         x->right = T2;
98
99         update_height(x);
100        update_height(y);
101
102        return y;
103    }
104
105    int balance_factor(Node* node) const {
106        return node ? height(node->left) - height(node->right) : 0;
107    }
108
109    Node* balance(Node* node) {
110        update_height(node);
111        int bf = balance_factor(node);
112
113        // Left Heavy
114        if (bf > 1) {
115            if (balance_factor(node->left) < 0)
116                node->left = rotate_left(node->left);
117            return rotate_right(node);
118        }
119        // Right Heavy
120        if (bf < -1) {
121            if (balance_factor(node->right) > 0)
122                node->right = rotate_right(node->right);
123            return rotate_left(node);
124        }
125        return node;

```

```

126     }
127
128     template <typename Func>
129     void traverse_in_order(Node* node, Func f) const {
130         if (!node) return;
131         std::stack<Node*> stack;
132         Node* current = node;
133         while (current || !stack.empty()) {
134             while (current) {
135                 stack.push(current);
136                 current = current->left;
137             }
138             current = stack.top();
139             stack.pop();
140             f(current->key);
141             current = current->right;
142         }
143     }
144
145     Node* insert(Node* node, const T& key) {
146         if (!node) {
147             return create_node(key);
148         }
149
150         if (comp_(key, node->key)) {
151             node->left = insert(node->left, key);
152         } else if (comp_(node->key, key)) {
153             node->right = insert(node->right, key);
154         } else {
155             return node;
156         }
157
158         return balance(node);
159     }
160
161 public:
162     AVLSet(Compare Comp = Compare()) : root(nullptr), size(0), comp_(Comp) {}
163     Compare comparator() const { return comp_; }
164     // Move operations
165     AVLSet(AVLSet&& other) noexcept
166         : root(other.root),
167         size(other.size){
168         other.root = nullptr;
169         other.size = 0;
170     }
171
172     AVLSet& operator=(AVLSet&& other) noexcept {
173         if (this != &other) {
174             clear();
175             root = other.root;
176             size = other.size;
177             other.root = nullptr;
178             other.size = 0;
179         }
180         return *this;
181     }
182

```

```

183 // Disable copy operations
184 AVLSet(const AVLSet&) = delete;
185 AVLSet& operator=(const AVLSet&) = delete;
186
187 void clear() {
188     delete_tree(root);
189     root = nullptr;
190 }
191
192 bool empty() const {
193     return size == 0;
194 }
195
196 size_t get_size() const {
197     return size;
198 }
199
200 template <typename Func>
201 void traverse_in_order(Func f) const {
202     traverse_in_order(root, f);
203 }
204
205 std::vector<T> items() const {
206     std::vector<T> result;
207     traverse_in_order([&result](const T& key) {
208         result.push_back(key);
209     });
210     return result;
211 }
212
213
214 void swap_with(AVLSet& other) {
215     std::swap(root, other.root);
216     std::swap(size, other.size);
217 }
218
219 void insert(const T& val){
220     root = insert(root, val);
221 }
222 void remove(const T& val){
223     root = remove(root, val);
224 }
225
226 template<typename RandAccIt>
227 void construct(RandAccIt bg, RandAccIt ed){
228     // Clear the current tree
229     delete_tree(root);
230     root = nullptr;
231     // Build new tree
232     root = build(bg, ed);
233     size = static_cast<size_t>(ed - bg);
234 }
235
236 /** Merge two sets in O(N+M) time.
237  * Some additional space may be costed.
238  * But it does not affect the result of the experiment.
239  * @param other The AVLSet to be merged into this set.

```

```

240  */
241
242  void linearmerge(AVLSet&& other) {
243      if (other.empty()) return;
244
245      // Get elements from both trees
246      std::vector<T> q1 = items();
247      std::vector<T> q2 = other.items();
248      std::vector<T> all_elements;
249      all_elements.reserve(q1.size() + q2.size());
250
251      // Merge sorted vectors
252      std::merge(q1.begin(), q1.end(), q2.begin(), q2.end(),
253                std::back_inserter(all_elements), comp_);
254
255      // Recycle both trees
256      delete_tree(root);
257      delete_tree(other.root);
258      root = nullptr;
259      other.root = nullptr;
260
261      // Build new tree
262      root = build(all_elements.begin(), all_elements.end());
263      size = all_elements.size();
264  }
265
266  /**
267   * Merge two sets in  $O(N \log(M))$  time.
268   * @param other The AVLSet to be merged into this set.
269   */
270  void simplemerge(AVLSet&& other) {
271      if (other.empty()) return;
272
273      if (size < other.size) { swap_with(other); }
274
275      // Insert all elements from the smaller tree (now 'other') into this
276      other.traverse_in_order([this](const T& key) {
277          this->insert(key);
278      });
279
280      other.clear();
281  }
282
283  void brownmerge(AVLSet&& other) {
284      if (other.empty()) return;
285      if (size < other.size) swap_with(other);
286
287      std::vector<T> elems = other.items();
288      other.clear();
289
290      // stacks of pointers-to-links (Node**). Each points to some parent->left or
291      // parent->right or &root
292      std::vector<Node**> path;
293      std::vector<Node**> successor;
294
295      path.push_back(&root);

```



```

296     for (const T& x : elems) {
297         // ==== CLIMB/RETRACT: pop successors while their key <= x (i.e. not (x <
298             succ_key)) ====
299         while (!successor.empty() && !comp_(x, (*successor.back())->key)) {
300             Node** succLink = successor.back();
301             // pop path entries until top == succLink
302             while (!path.empty() && path.back() != succLink) path.pop_back();
303             successor.pop_back();
304         }
305
306         // ==== DESCEND from current finger (path.back()) to insertion point ====
307         Node** curLink = path.empty() ? &root : path.back();
308         Node* p = *curLink;
309         if (!p) {
310             // empty subtree (rare because root existed), insert directly
311             *curLink = create_node(x);
312             path.push_back(curLink);
313         } else {
314             for (;;) {
315                 path.push_back(curLink); // link that points to p
316                 if (comp_(x, p->key)) {
317                     // go left; mark this link as a successor (we turned left here)
318                     if (p->left == nullptr) {
319                         p->left = create_node(x);
320                         path.push_back(&(p->left));
321                         break;
322                     } else {
323                         successor.push_back(curLink);
324                         curLink = &(p->left);
325                         p = *curLink;
326                     }
327                 } else {
328                     // go right
329                     if (p->right == nullptr) {
330                         p->right = create_node(x);
331                         path.push_back(&(p->right));
332                         break;
333                     } else {
334                         curLink = &(p->right);
335                         p = *curLink;
336                     }
337                 }
338             }
339         }
340
341         // ==== REBALANCE upward using the link-stack; attach rotated subtree via *
342             link ====
343         while (!path.empty()) {
344             Node** link = path.back();
345             Node* s = *link;
346             path.pop_back();
347
348             if (!successor.empty() && successor.back() == link) successor.pop_back();
349
350             update_height(s);
351             int bf = balance_factor(s);

```

```

351         if (std::abs(bf) > 1) {
352             Node* newsub = balance(s); // returns new root of this subtree
353             *link = newsub; // reattach correctly via the link
354
355             // retract path until the remaining entries are consistent with this
356             // rotation
357             while (!path.empty() && path.back() != link) path.pop_back();
358             break; // stop climbing after performing rotation
359         }
360
361         if (bf == 0) {
362             // height didn't increase => stop climbing
363             break;
364         }
365         // else continue climbing
366     }
367     #ifdef DEBUG
368     auto items_after = items();
369     for (size_t i = 1; i < items_after.size(); ++i)
370         assert(!comp_(items_after[i], items_after[i-1]));
371     #endif
372 }
373
374 };

```

References

- Brown, M. R., & Tarjan, R. E. (1979). A fast merging algorithm. *Journal of the ACM (JACM)*, 26(2), 211–226.
- Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2022). *Introduction to algorithms* (Fourth ed.). MIT Press.
- Karltun, P. L., Fuller, S. H., Scroggs, R., & Kaehler, E. (1976). Performance of height-balanced trees. *Communications of the ACM*, 19(1), 23–28.
- Stockmeyer, P. K., & Yao, F. F. (1980). On the optimality of linear merge. *SIAM Journal on Computing*, 9(1), 85–90.